
module *Reachable*

This module specifies an algorithm for computing the set of nodes in a directed graph that are reachable from a given node called *Root*. The algorithm is due to *Jayadev Misra*. It is, to my knowledge, a new variant of a fairly obvious breadth-first search for reachable nodes. I find this algorithm interesting because it is easier to implement using multiple processors than the obvious algorithm. Module *ParReach* describes such an implementation. You may want to read it after reading this module.

Module *ReachableProofs* contains a TLA+ proof of the algorithm's safety property—that is, partial correctness, which means that if the algorithm terminates then it produces the correct answer. That proof has been checked by *TLAPS*, the TLA+ proof system. The proof is based on ideas from an informal correctness proof by *Misra*.

In this module, reachability is expressed in terms of the operator *ReachableFrom*, where *ReachableFrom(S)* is the set of nodes reachable from the nodes in the set *S* of nodes. This operator is defined in module *Reachability*. That module describes a directed graph in terms of the constants *Nodes* and *Succ*, where *Nodes* is the set of nodes and *Succ* is a function with domain *Nodes* such that *Succ[m]* is the set of all nodes *n* such that there is an edge from *m* to *n*. If you are not familiar with directed graphs, you should read at least the opening comments in module *Reachability*.

extends *Reachability*; *Integers*; *FiniteSets*

constant *Root*

assume *RootAssump* \triangleq *Root* \in *Nodes*

Reachable is defined to be the set of nodes reachable from *Root*. The purpose of the algorithm is to compute *Reachable*.

Reachable \triangleq *ReachableFrom*({*Root*})

The obvious algorithm for computing *Reachable*({*Root*}) is as follows. There are two variables which, following *Misra*, we call *marked* and *vroot*. Each variable holds a set of nodes that are reachable from *Root*. Initially, *marked* = {} and *vroot* = {*Root*}. While *vroot* is non-empty, the obvious algorithm removed an arbitrary node *v* from *vroot*, adds *v* to *marked*, and adds to *vroot* all nodes in *Succ[v]* that are not in *marked*. The algorithm terminates when *vroot* is empty, which will eventually be the case if and only if *Reachable*({*Root*}) is a finite set. When it terminates, *marked* equals *Reachable*({*Root*}).

In the obvious algorithm, *marked* and *vroot* are always disjoint sets of nodes. *Misra*'s variant differs in that *marked* and *vroot* are not necessarily disjoint. While *vroot* is nonempty, it chooses an arbitrary node and does the following:

```

if v is not in in marked then it performs the same action as the obvious algorithm except:
    (1) it doesn't remove v from vroot, and
    (2) it adds all nodes in Succ[v] to vroot, not just the ones not in marked .
else it removes v from vroot

```

Here is the algorithm's *PlusCal* code.

```

--fair algorithm Reachable f
variables marked = {}; vroot = {Root};
f a: while ( vroot  $\neq$  {} )
    f with ( v  $\in$  vroot )
        f if ( v  $\in$  marked )
            f marked := marked  $\cup$  {v};

```

```

          vroot := vroot ∪ Succ[v] g
        else f vroot := vroot \ {v} g
      g
    g
  g
g

```

Here is the TLA+ translation of the *PlusCal* code.

BEGIN TRANSLATION

variables *marked*; *vroot*; *pc*

Reachable \triangleq *ReachableFrom*(*marked*)

vars \triangleq \langle *marked*; *vroot*; *pc* \rangle

Init \triangleq Global variables
 \wedge *marked* = {}
 \wedge *vroot* = {*Root*}
 \wedge *pc* = "a"

a \triangleq \wedge *pc* = "a"
 \wedge if *vroot* \neq {}
 then $\wedge \exists v \in$ *vroot* :
 if *v* \in *marked*
 then \wedge *marked*' = (*marked* \cup {*v*})
 \wedge *vroot*' = (*vroot* \cup *Succ*[*v*])
 else \wedge *vroot*' = *vroot* \ {*v*}
 \wedge unchanged *marked*
 \wedge *pc*' = "a"
 else \wedge *pc*' = "Done"
 \wedge unchanged \langle *marked*; *vroot* \rangle

Next \triangleq *a*
 \vee Disjunct to prevent deadlock on termination
 (*pc* = "Done" \wedge unchanged *vars*)

Spec \triangleq \wedge *Init* $\wedge \Box$ [*Next*]_{*vars*}
 \wedge WF_{*vars*}(*Next*)

Termination \triangleq \Diamond (*pc* = "Done")

END TRANSLATION

Partial correctness is based on the invariance of the following four state predicates. I have sketched very informal proofs of their invariance, as well of proofs of the the two theorems that assert correctness of the algorithm. The module *ReachableProofs* contains rigorous, *TLAPS* checked TLA+ proofs of all except the last theorem. The last theorem asserts termination, which is a liveness property, and *TLAPS* is not yet capable of proving liveness properties.

$$\begin{aligned}
TypeOK &\triangleq \wedge marked \in \text{subset } Nodes \\
&\wedge vroot \in \text{subset } Nodes \\
&\wedge pc \in \{ "a"; "Done" \} \\
&\wedge (pc = "Done") \Rightarrow (vroot = \{\})
\end{aligned}$$

The invariance of *TypeOK* is obvious. (I decided to make the obvious fact that *pc* equals *\Done* only if *vroot* is empty part of the type-correctness invariant.)

$$\begin{aligned}
Inv1 &\triangleq \wedge TypeOK \\
&\wedge \forall n \in marked : Succ[n] \subseteq (marked \cup vroot)
\end{aligned}$$

The second conjunct of *Inv1* is invariant because each element of *Succ*[*n*] is added to *vroot* when *n* is added to *marked*, and it remains in *vroot* at least until it's added to *marked*. I made *TypeOK* a conjunct of *Inv1* to make *Inv1* an inductive invariant, which made the TLA+ proof of its invariance a tiny bit easier to read.

$$Inv2 \triangleq (marked \cup ReachableFrom(vroot)) = ReachableFrom(marked \cup vroot)$$

Since *ReachableFrom*(*marked* \cup *vroot*) is the union of *ReachableFrom*(*marked*) and *ReachableFrom*(*vroot*), to prove that *Inv2* is invariant we must show *ReachableFrom*(*marked*) is a subset of *marked* \cup *ReachableFrom*(*vroot*). For this, we assume that *m* is in *ReachableFrom*(*marked*) and show that it either is in *marked* or is reachable from a node in *vroot*.

Since *m* is in *ReachableFrom*(*marked*), there is a path with nodes *p*₁, *p*₂, ..., *p*_j such that *p*₁ is in *marked* and *p*_j = *m*. If all the *p*_i are in *marked*, then *m* is in *marked* and we're done. Otherwise, choose *i* such that *p*₁, ..., *p*_i are in *marked*, but *p*_(i+1) isn't in *marked*. Then *p*_(i+1) is in *succ*[*p*_i], which by *Inv1* implies that it's in *marked* \cup *vroot*. Since it isn't in *marked*, it must be in *vroot*. The path with nodes *p*_(i+1), ..., *p*_j shows that *p*_j, which equals *m*, is in *ReachableFrom*(*vroot*). This completes the proof that *m* is in *marked* or *ReachableFrom*(*vroot*).

$$Inv3 \triangleq Reachable = marked \cup ReachableFrom(vroot)$$

For convenience, let *R* equal *marked* \cup *ReachableFrom*(*vroot*). In the initial state, *marked* = {} and *vroot* = {*Root*}, so *R* equals *Reachable* and *Inv3* is true. We have to show that each action *a* step leaves *R* unchanged. There are two cases:

Case1: The *a* step adds an element *v* of *vroot* to *marked* and adds to *vroot* the nodes in *Succ*[*v*], which are all in *ReachableFrom*(*vroot*). Since *v* itself is also in *ReachableFrom*(*vroot*), the step leaves *R* unchanged.

Case 2: The *a* step removes from *vroot* an element *v* of *marked*. Since *Inv1* implies that every node in *Succ*[*v*] is in *vroot*, the only element that this step removes from *ReachableFrom*(*vroot*) is *v*, which the step adds to *marked*. Hence *R* is unchanged.

It is straightforward to use *TLC* to check that *Inv1* – *Inv3* are invariants of the algorithm for small graphs.

Partial correctness of the algorithm means that if it has terminated, then *marked* equals *Reachable*. The algorithm has terminated when *pc* equals *\Done*, so this theorem asserts partial correctness.

theorem *Spec* $\Rightarrow \Box((pc = "Done") \Rightarrow (marked = Reachable))$

TypeOK implies $(pc = \backslash Done) \Rightarrow (vroot = \{\})$. Since, *ReachableFrom*({}) equals {}, *Inv3* implies $(vroot = \{\}) \Rightarrow (marked = Reachable)$. Hence the theorem follows from the invariance of *TypeOK* and *Inv3*.

The following theorem asserts that if the set of nodes reachable from *Root* is finite, then the algorithm eventually terminates. Of course, this liveness property can be true only because *Spec* implies weak fairness of *Next*, which equals action *a* .

```
theorem assume IsFiniteSet(Reachable)
  prove Spec  $\Rightarrow$   $\Diamond(pc = \text{"Done"})$ 
```

To prove the theorem, we assume a behavior satisfies *Spec* and prove that it satisfies $\Diamond(pc = \text{"Done"})$. If *pc* = *\a* and *vroot* = {}, then an *a* step sets *pc* to *\Done*. Since invariance of *TypeOK* implies $\Box(pc \in \{\a; \text{"Done"}\})$, weak fairness of *a* implies that to prove $\Diamond(pc = \text{"Done"})$, it suffices to prove $\Diamond(vroot = \{\})$.

We prove $\Diamond(vroot = \{\})$ by contradiction. We assume it's false, which means that $\Box(vroot \neq \{\})$ is true, and obtain a contradiction. From

$\Box \textit{TypeOK}$, we infer that $\Box(vroot \neq \{\})$ implies $\Box(pc = \a)$. By weak fairness of action *a* , $\Box(vroot \neq \{\})$ implies that there are an infinite number of *a* steps. The assumption that *Reachable* is finite and $\Box \textit{Inv3}$ imply that *marked* and *vroot* are always finite. Since *vroot* is always finite and nonempty, from any state there can be only a finite number of *a* steps that remove an element from *vroot* until there is an *a* step that adds a new element to *marked* . Since there are an infinite number of *a* steps, there must be an infinite number of steps that add new elements to *marked* . This is impossible because *marked* is a finite set. Hence, we have the required contradiction.

TLC can quickly check these two theorems on models containing a half dozen nodes.

```
\ * Modification History
\ * Last modified Wed Apr 17 12:21:22 PDT 2019 by lamport
\ * Created Thu Apr 04 10:11:51 PDT 2019 by lamport
```